

# Jug Tutorial: Coarse-Level Parallel Programming in Python

Luis Pedro Coelho

[lpc@cmu.edu](mailto:lpc@cmu.edu)

January 29, 2010



Carnegie Mellon

# Problem

Brute force a password.

# Assumptions

The **crypt** module exists with elements:

- decrypt: decrypt ciphertext given a password.
- isgood: test whether this text could be the plaintext.
- letters: just the letters.

Also, we know that the password is five letters.

# Simple Solution

```
import itertools
from crypt import decode, letters, isgood, preprocess

ciphertext = file('secret.msg').read()
ciphertext = preprocess(ciphertext)

for p in itertools.product(letters, repeat=5):
    text = decode(ciphertext, p)
    if isgood(text):
        passwd = "".join(map(chr, p))
        print '%s:%s' % (passwd, text)
```

This solution does not take advantage of multiple processors.

# Tasks

```
import itertools
from crypt import decode, letters, isgood, preprocess

ciphertext = file('secret.msg').read()
ciphertext = preprocess(ciphertext)

def decrypt(ciphertext, p):
    text = decode(ciphertext, p)
    if isgood(text):
        passwd = "".join(map(chr, p))
        return (passwd, text)
    # else: return None

results = []
for p in itertools.product(letters, repeat=5):
    results.append(Task(decrypt, ciphertext, p))
```

# Python Magic

```
from jug import TaskGenerator
import itertools
from crypt import decode, letters, isgood, preprocess

ciphertext = file('secret.msg').read()
ciphertext = preprocess(ciphertext)

@TaskGenerator
def decrypt(ciphertext, p):
    text = decode(ciphertext, p)
    if isgood(text):
        passwd = "".join(map(chr, p))
        return (passwd, text)
    # else: return None

results = []
for p in itertools.product(letters, repeat=5):
    results.append(decrypt(ciphertext, p))
```

# Enter Jug

You give it the Jugfile, it runs the tasks for you!



# Jug Loop

```
while len(tasks) > 0:
    ready = [t for t in tasks if can_run(t)]
    for t in ready:
        if not is_running(t):
            t.run()
        tasks.remove(t)
```

Except jug is much fancier!

# Jug Advantages

- ➊ Automatic task-level parallelization with dependency tracking.
- ➋ Remember all intermediate results.
- ➌ Makes writing parallel code look like writing sequential code.

This example is actually not so good.

We have  $26^5 \approx 11M$  tasks, all of which run very fast.

As a rule of thumb, your tasks should take at least a couple of seconds.

**Solution** each task will be:

Given a letter, try **all passwords** beginning with that letter.

Now, we have 26 tasks. Much better.

```

@TaskGenerator
def decrypt(prefix):
    res = []
    for suffix in product(letters, repeat=5-len(prefix)):
        passwd = np.concatenate([prefix, suffix])
        text = decode(ciphertext, passwd)
        if isgood(text):
            passwd = "".join(map(chr, passwd))
            res.append( (passwd, text) )
    return res

```

```

@TaskGenerator
def join(partials):
    return list(chain(*partials))

```

```

results = join([decrypt([p]) for p in letters])

```

Let's call **jug** now.

(Assuming our code is in a file called **jugfile.py**)

```
$jug status
```

| Task name       | Waiting | Ready | Finished | Running |
|-----------------|---------|-------|----------|---------|
| -----           |         |       |          |         |
| jugfile.join    | 1       | 0     | 0        | 0       |
| jugfile.decrypt | 0       | 26    | 0        | 0       |
| .....           |         |       |          |         |
| Total:          | 1       | 26    | 0        | 0       |

Some tasks are ready. None are running.

```
$jug execute &  
[1] 29501  
$jug execute &  
[2] 29502
```

Executing in the background...



```
$jug status
```

| Task name       | Waiting | Ready | Finished | Running |
|-----------------|---------|-------|----------|---------|
| -----           |         |       |          |         |
| jugfile.join    | 1       | 0     | 0        | 0       |
| jugfile.decrypt | 0       | 24    | 0        | 2       |
| .....           |         |       |          |         |
| Total:          | 1       | 24    | 0        | 2       |

Two tasks running. Good.

Wait a few minutes...

```
$jug status
```

| Task name       | Waiting | Ready | Finished | Running |
|-----------------|---------|-------|----------|---------|
| -----           |         |       |          |         |
| jugfile.join    | 1       | 0     | 0        | 0       |
| jugfile.decrypt | 0       | 14    | 10       | 2       |
| .....           |         |       |          |         |
| Total:          | 1       | 14    | 10       | 2       |

Ten tasks have finished.

Notice how the **join** task must wait for all the others.

A few more minutes. . .

```
$jug status
```

| Task name       | Waiting | Ready | Finished | Running |
|-----------------|---------|-------|----------|---------|
| jugfile.join    | 0       | 0     | 1        | 0       |
| jugfile.decrypt | 0       | 0     | 26       | 0       |
| .....           |         |       |          |         |
| Total:          | 0       | 0     | 27       | 0       |

# What Now?

All tasks are finished!  
How do I get to the results?

```
import jug
jug.init('jugfile', 'jugdata')
import jugfile
results = jug.task.value(jugfile.results)
for p,t in results:
    print "%s\n\n      Password was '%s'" % (t,p)
```

**jug.init** is necessary to initialise the jug backend.

Then, the **jugfile** can be directly imported as a Python module.

# For more information

<http://luispedro.org/software/jug>

<http://github.com/luispedro/jug>